

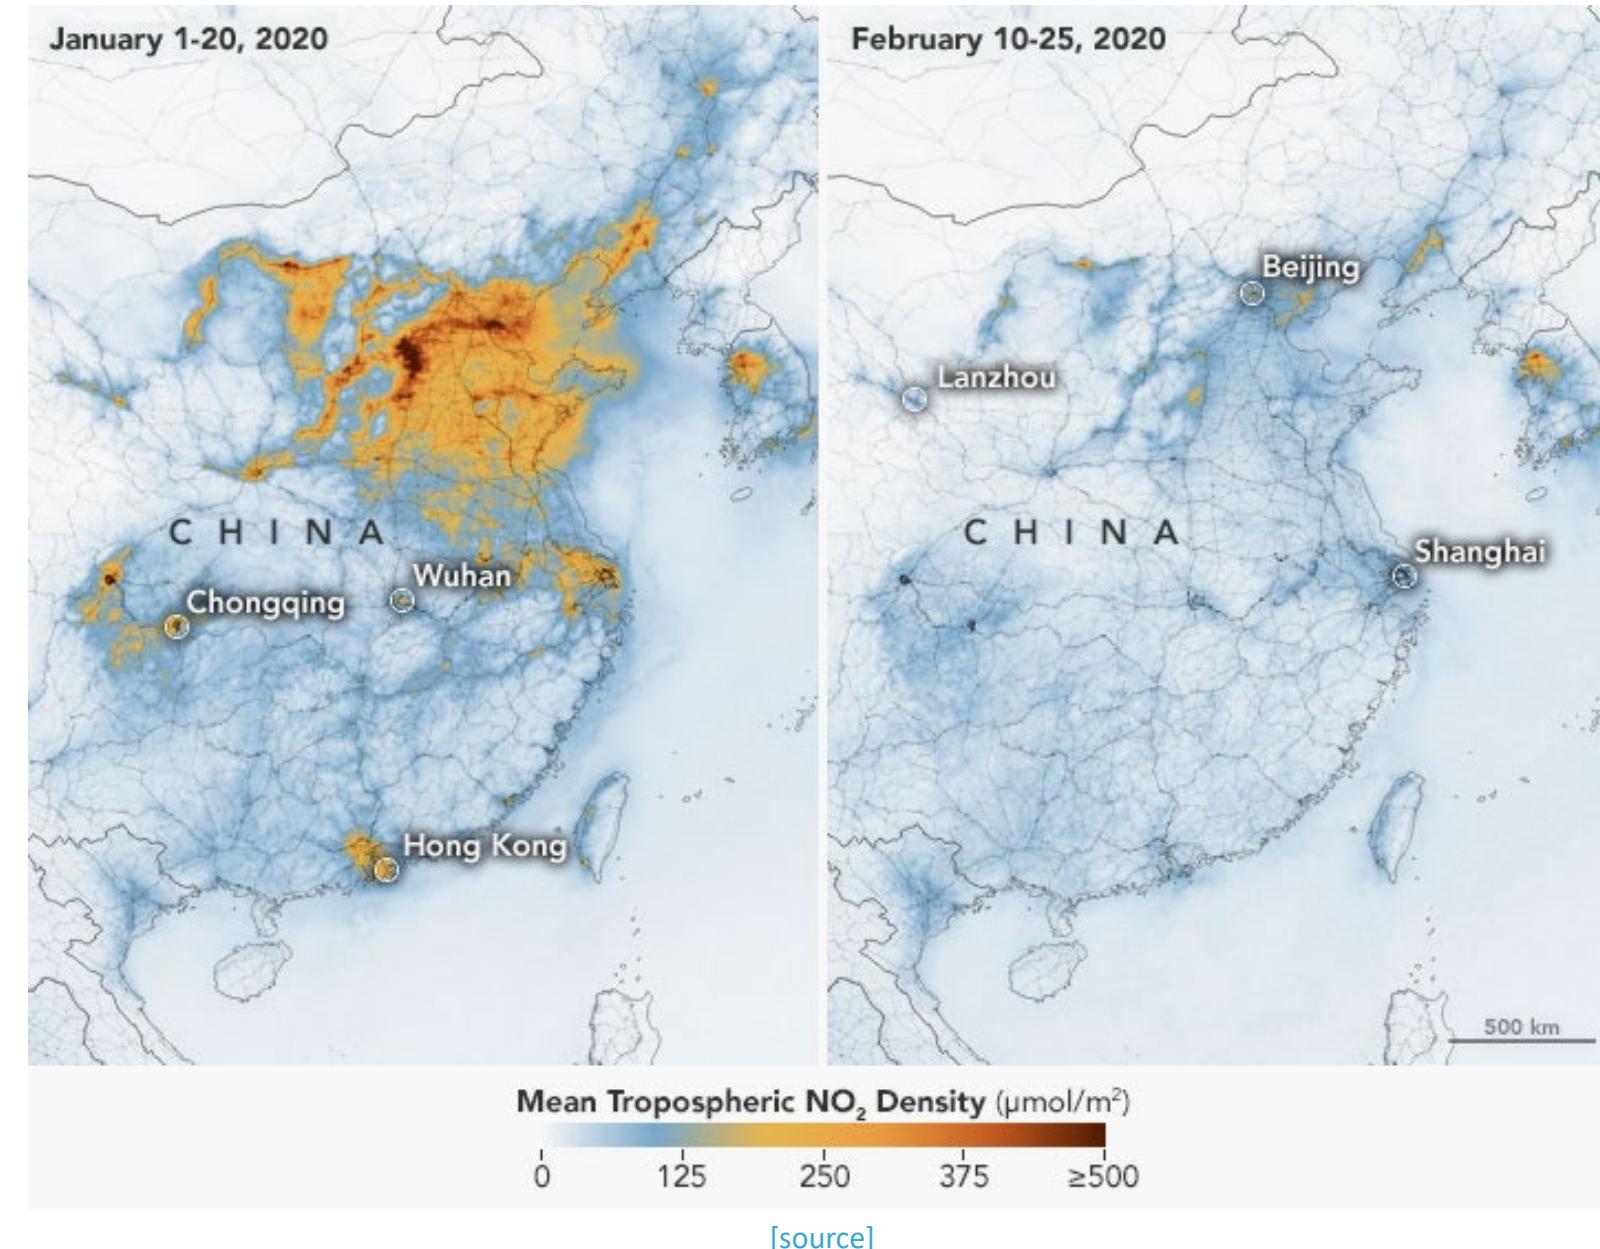
Introduction to *Urban* Data Science

Spatial Density Estimation

(EPA1316)

Lecture 13

Trivik Verma



Final Projects

A : Mobility, Built Environment & Sustainability

B : Identifying the Health Vulnerability in a City

C : Modelling COVID-19 in India

Final Projects - Milestones

1. Group Creation and Project Selection
2. Scope of Work and Preliminary EDA
 - Project statement
 - Preliminary EDA
3. EDA and Revised Project Statement
4. Project Report
 - Template
 - Rubric

Have fun!

Last Time

- Big Data and High Dimensionality
- A Framework For Dimensionality Reduction
- Principal Components Analysis (PCA)

Q: When does high-dimensionality occur?

- A. $P \gg N$
- B. $P \ll N$
- C. $P = N$
- D. When Probability of data is
Normally distributed

Q: What are some of the consequences of high-dimensional data?

- A. Issues with regression
- B. Variables related to each other
- C. Models cannot predict new data
- D. All the above
- E. None of the above

Today

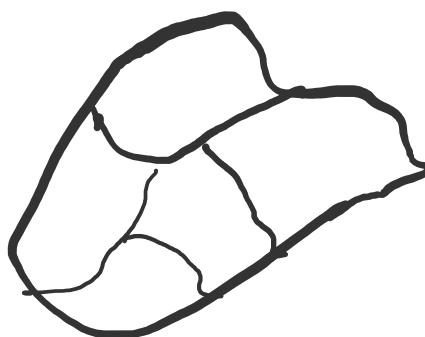
- The *point* of points
- Point patterns
- Visualization of point patterns
- Identifying clusters of points

The *point* of points

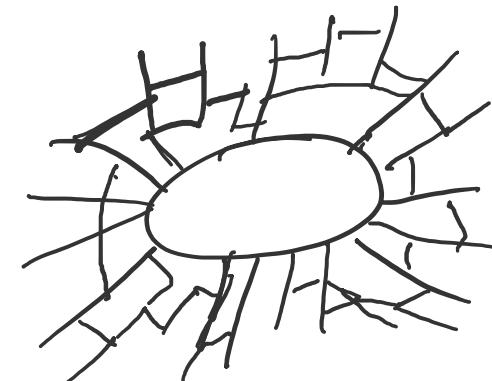
Points like polygons

- Points can represent “fixed” entities
- In this case, points are **qualitatively** like **polygons/lines**
- The **goal** here is, taking location fixed, to model other aspects of the data

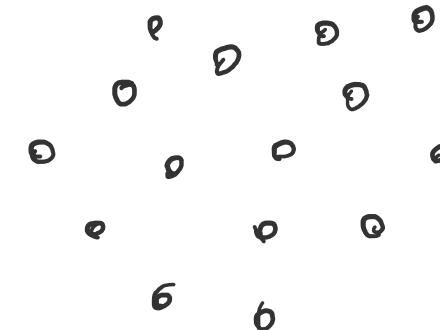
Polygons



Lines



Points



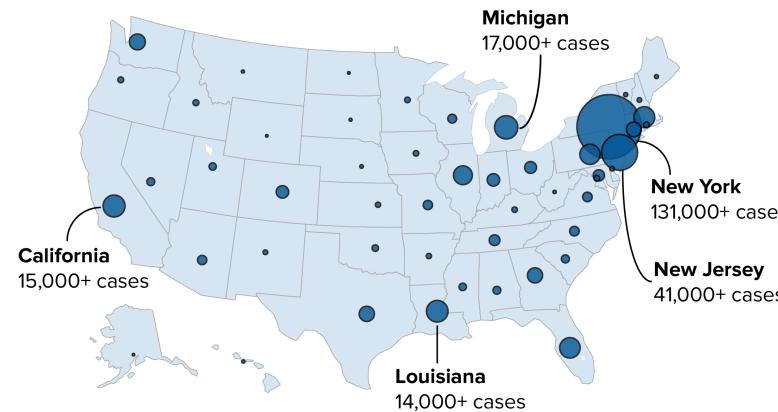
Points like polygons

Examples

- Cities (in most cases)
- Buildings or People (processes to estimate – social media)
- Polygons represented as their centroid ...

Reported coronavirus cases in the US

As of April 6, 2020



SOURCE: Johns Hopkins University. Data as of April 6, 2020 at 6 p.m. ET

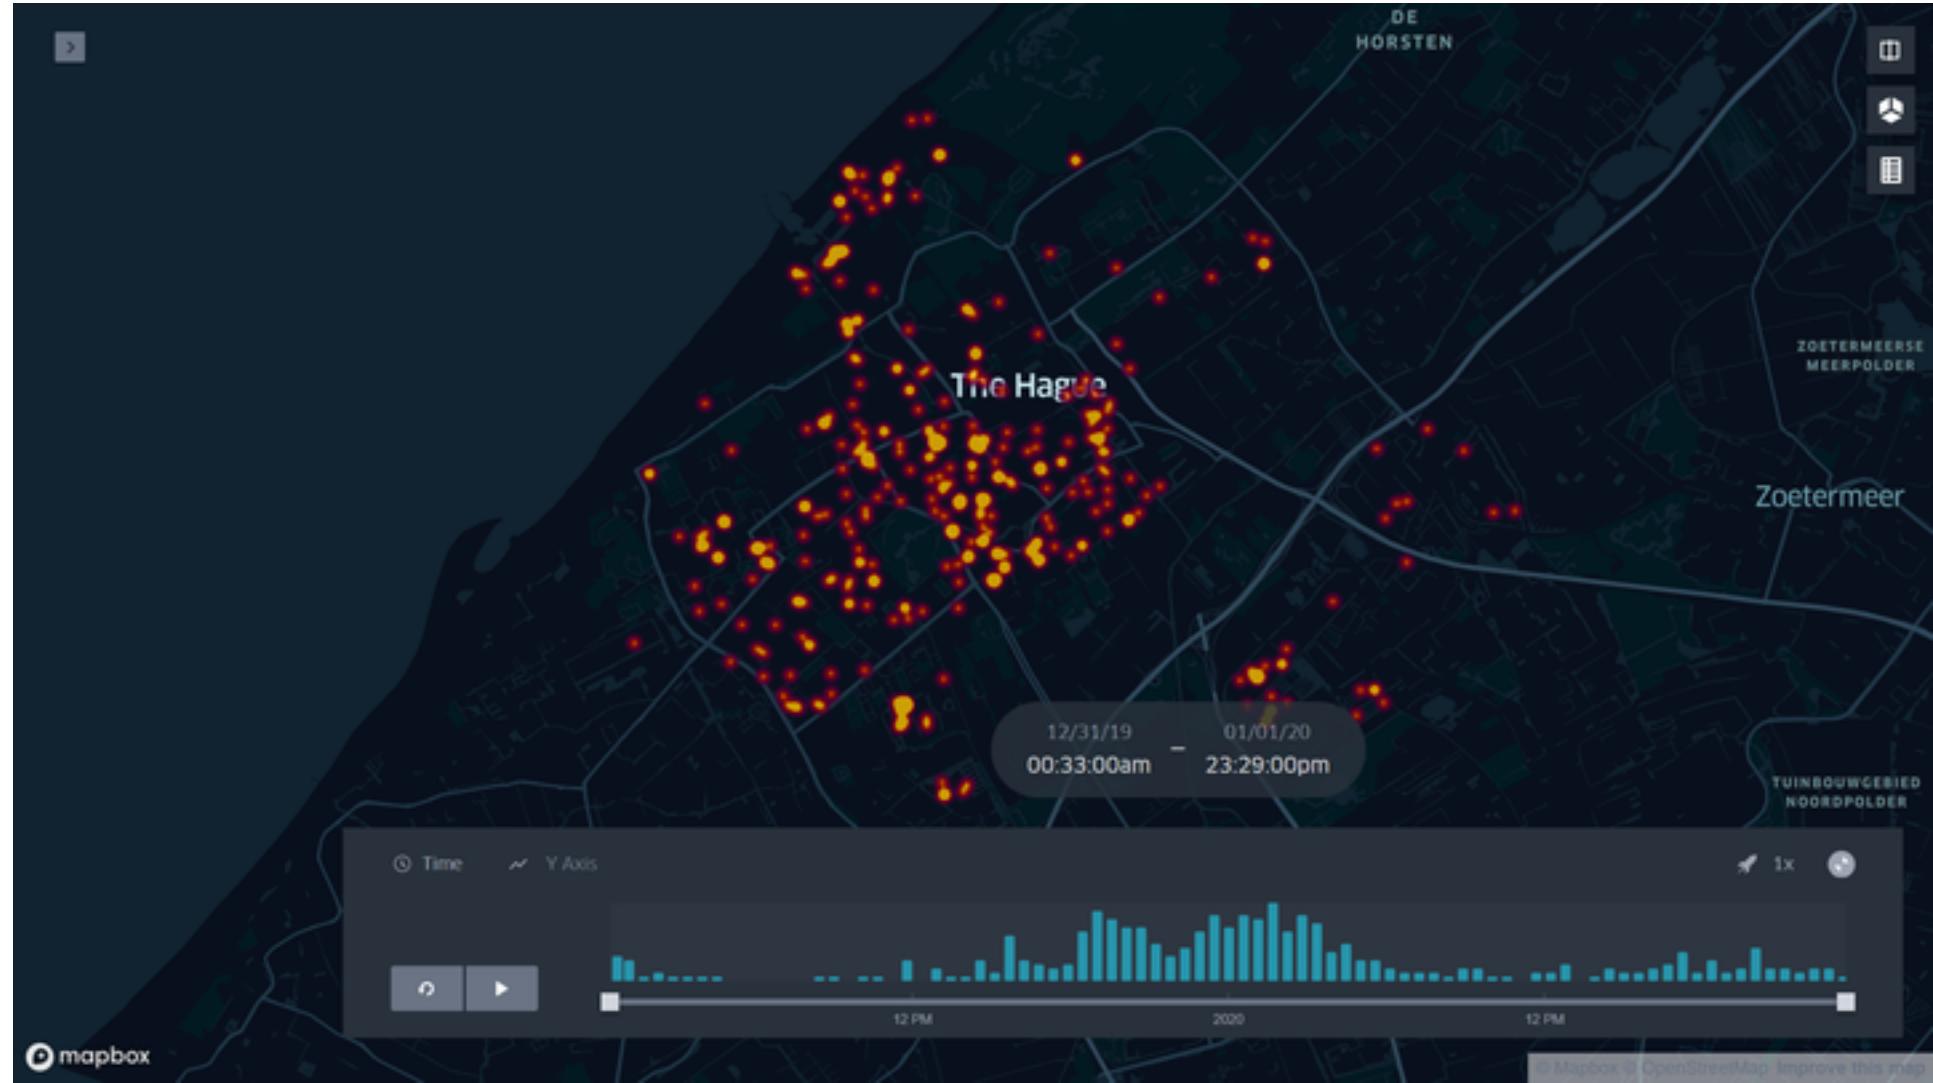


When points are not polygons

Point data are not only a different geometry than polygons or lines...

... Points can also represent a fundamentally different way to approach spatial analysis

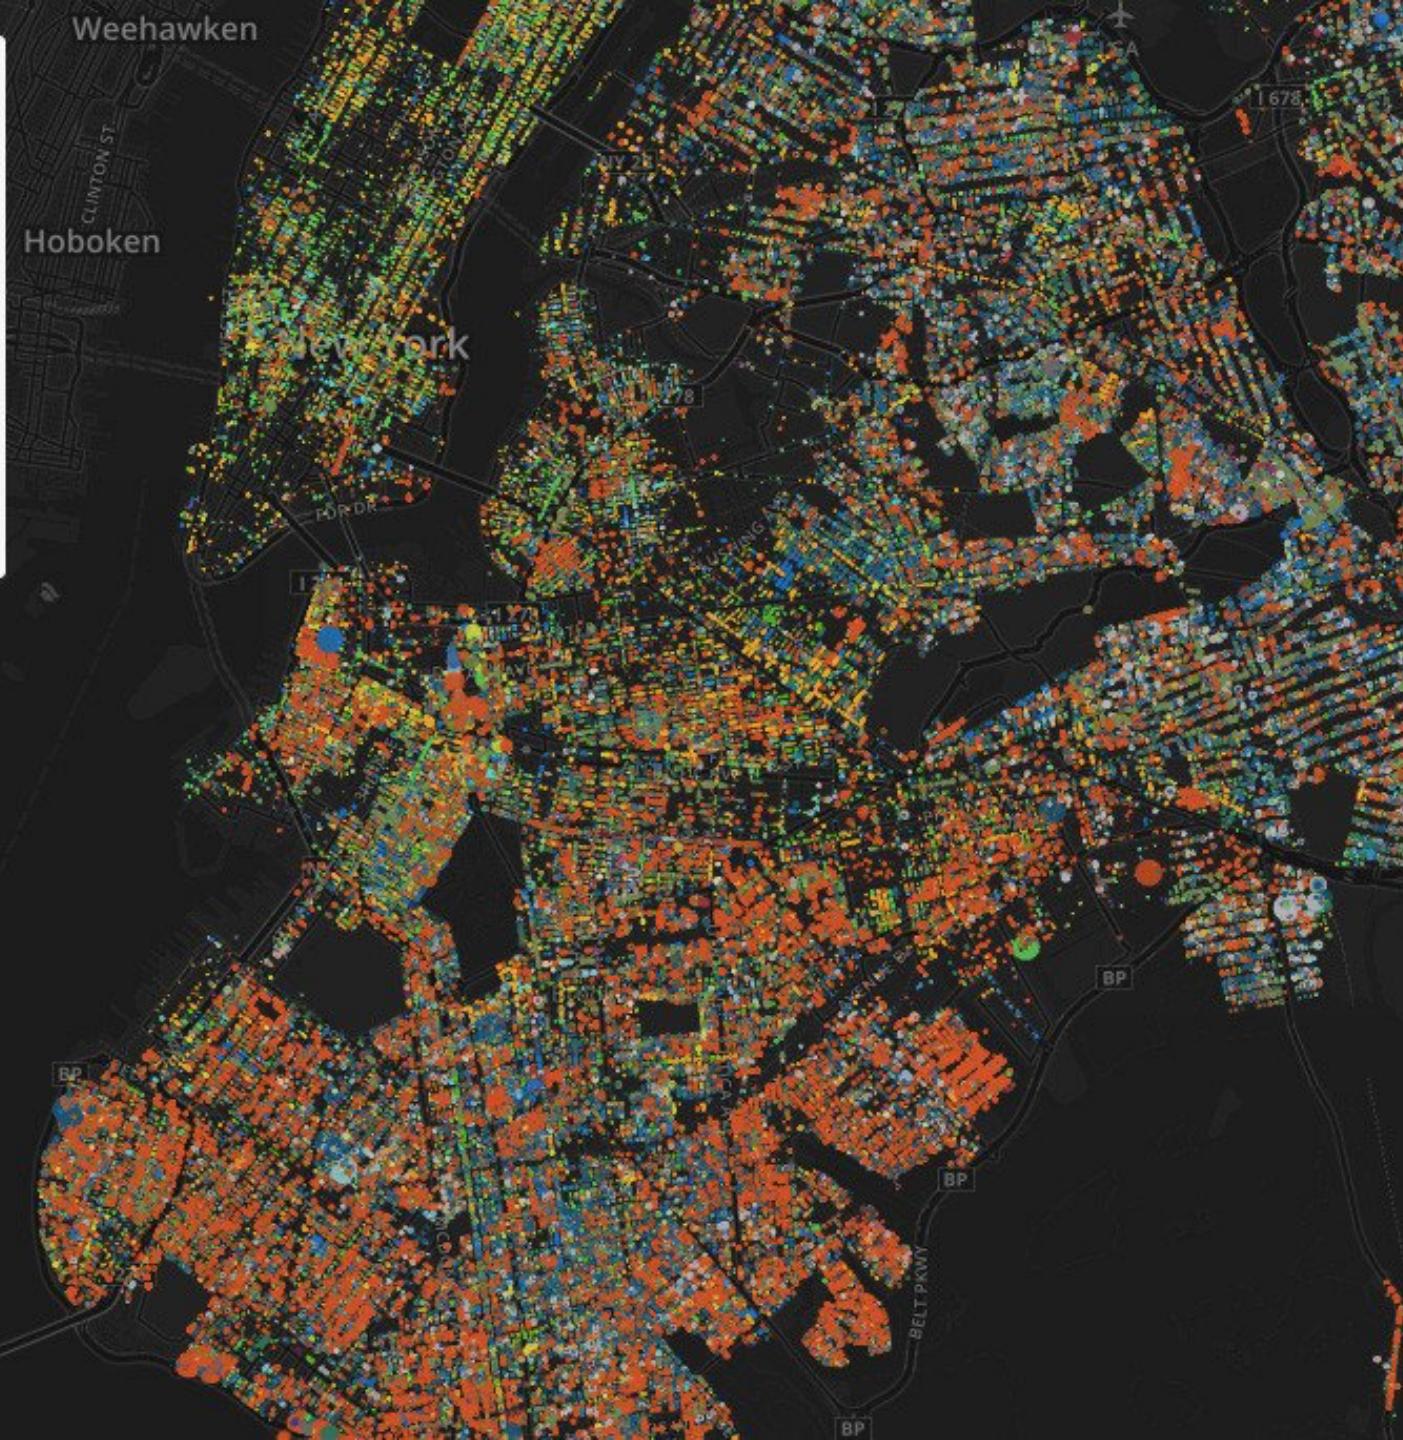
Points unlike polygons



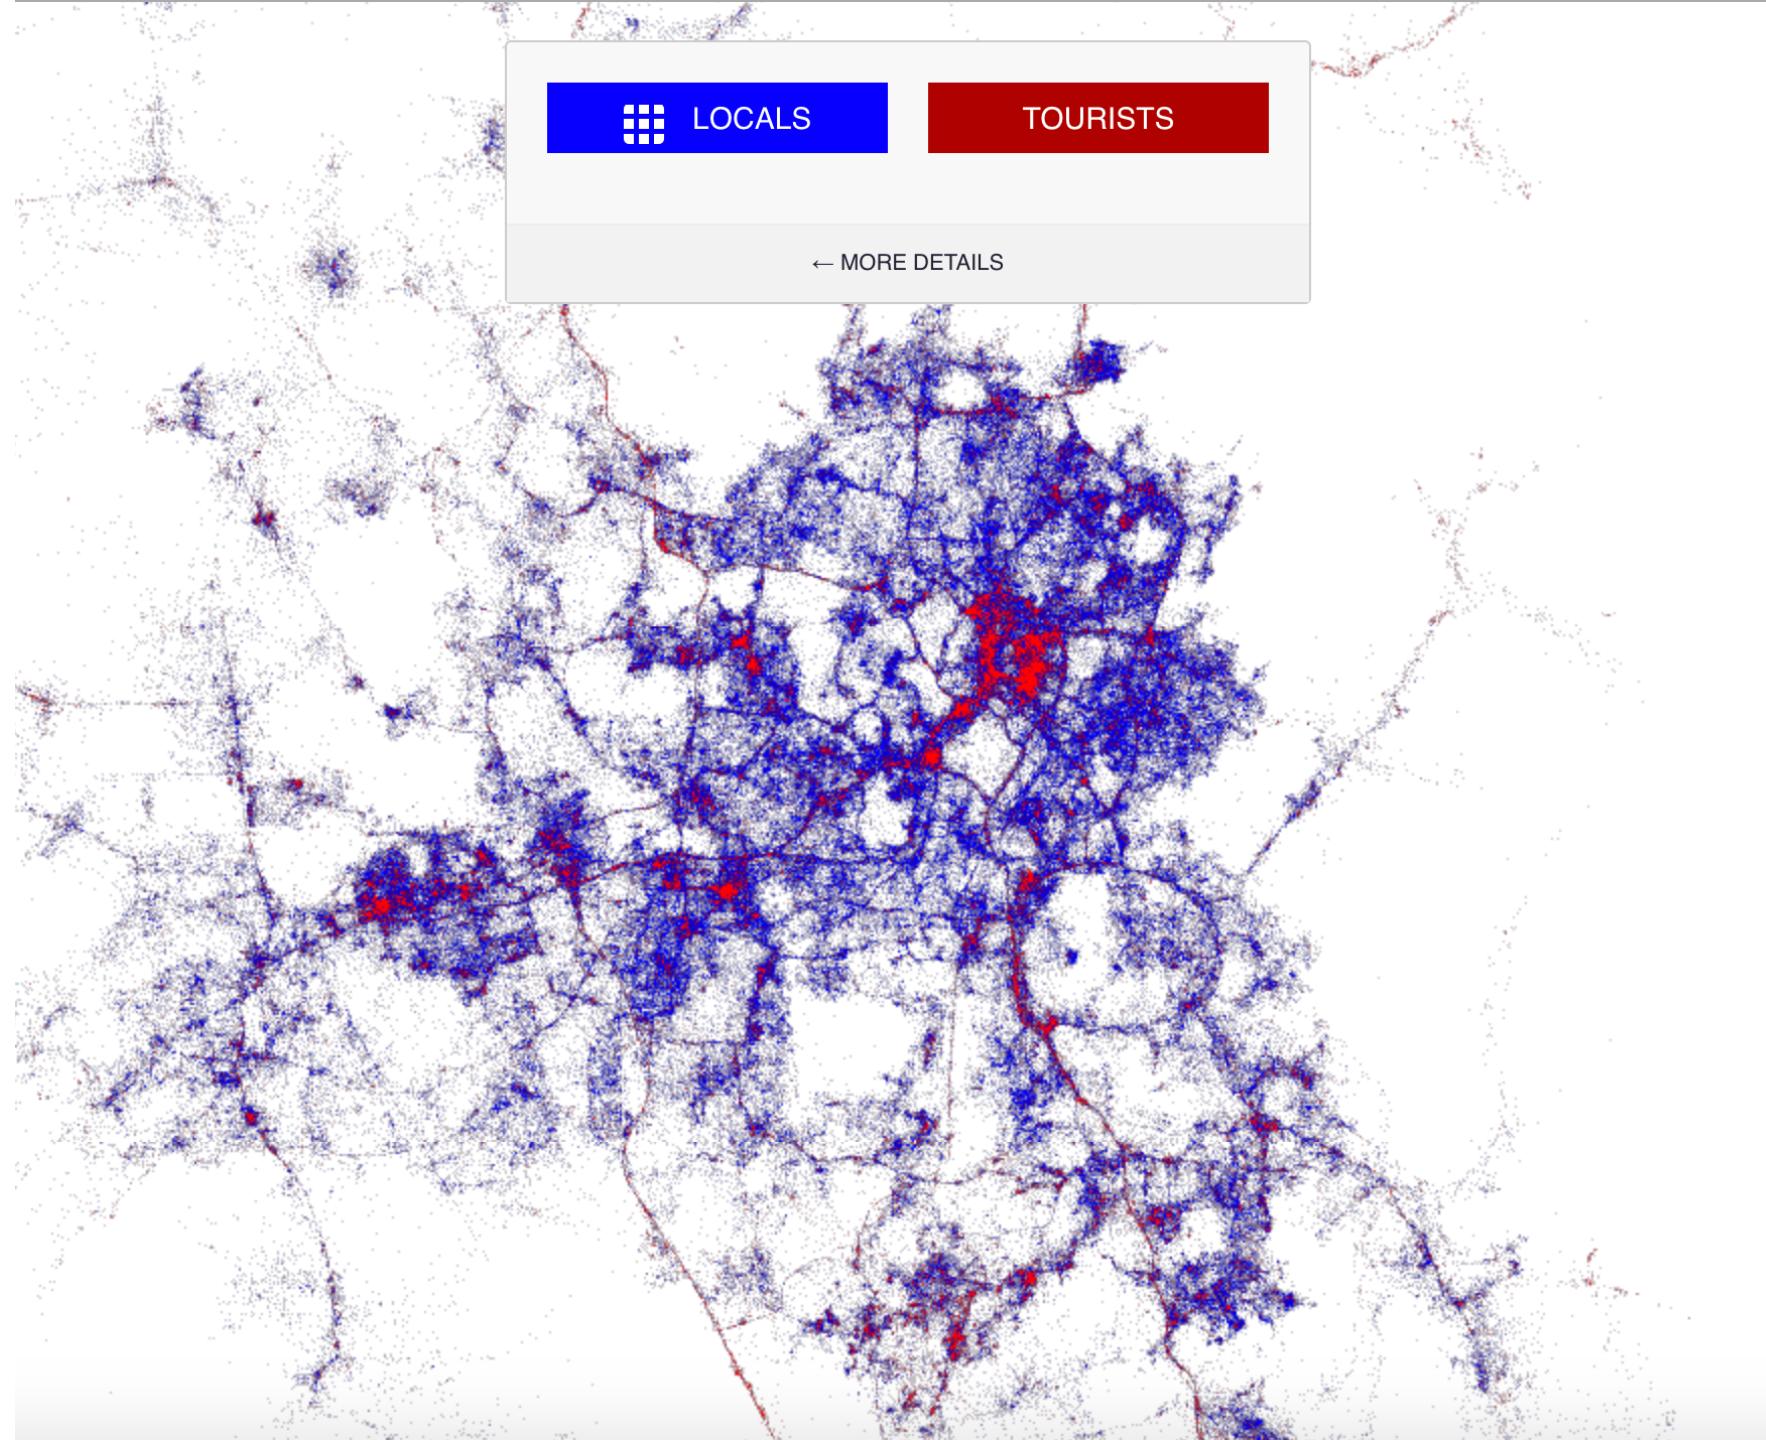
NYC Street Trees by Species

New York City's urban forest provides numerous environmental and social benefits, and street trees compose roughly one quarter of that canopy. This map shows the distribution and biodiversity of the city's street trees based on the last tree census.

[Read more.](#)



[\[source\]](#)



[source]

Point Patterns

Point Patterns

Distribution of **points over** a portion of space

Assumption is a point can happen anywhere on that space, but only happens in specific locations

- **Unmarked:** locations only
- **Marked:** values attached to each point

Point Pattern Analysis

Describe, characterize, and explain point patterns, focusing on their **generating process**

- Visual exploration
- Clustering properties and clusters
- Statistical modelling of the underlying processes

Visualisation of PPs

Visualisation of PPs

Two routes

1. *Aggregate* \leftrightarrow Histogram
2. *Smooth* \leftrightarrow KDE

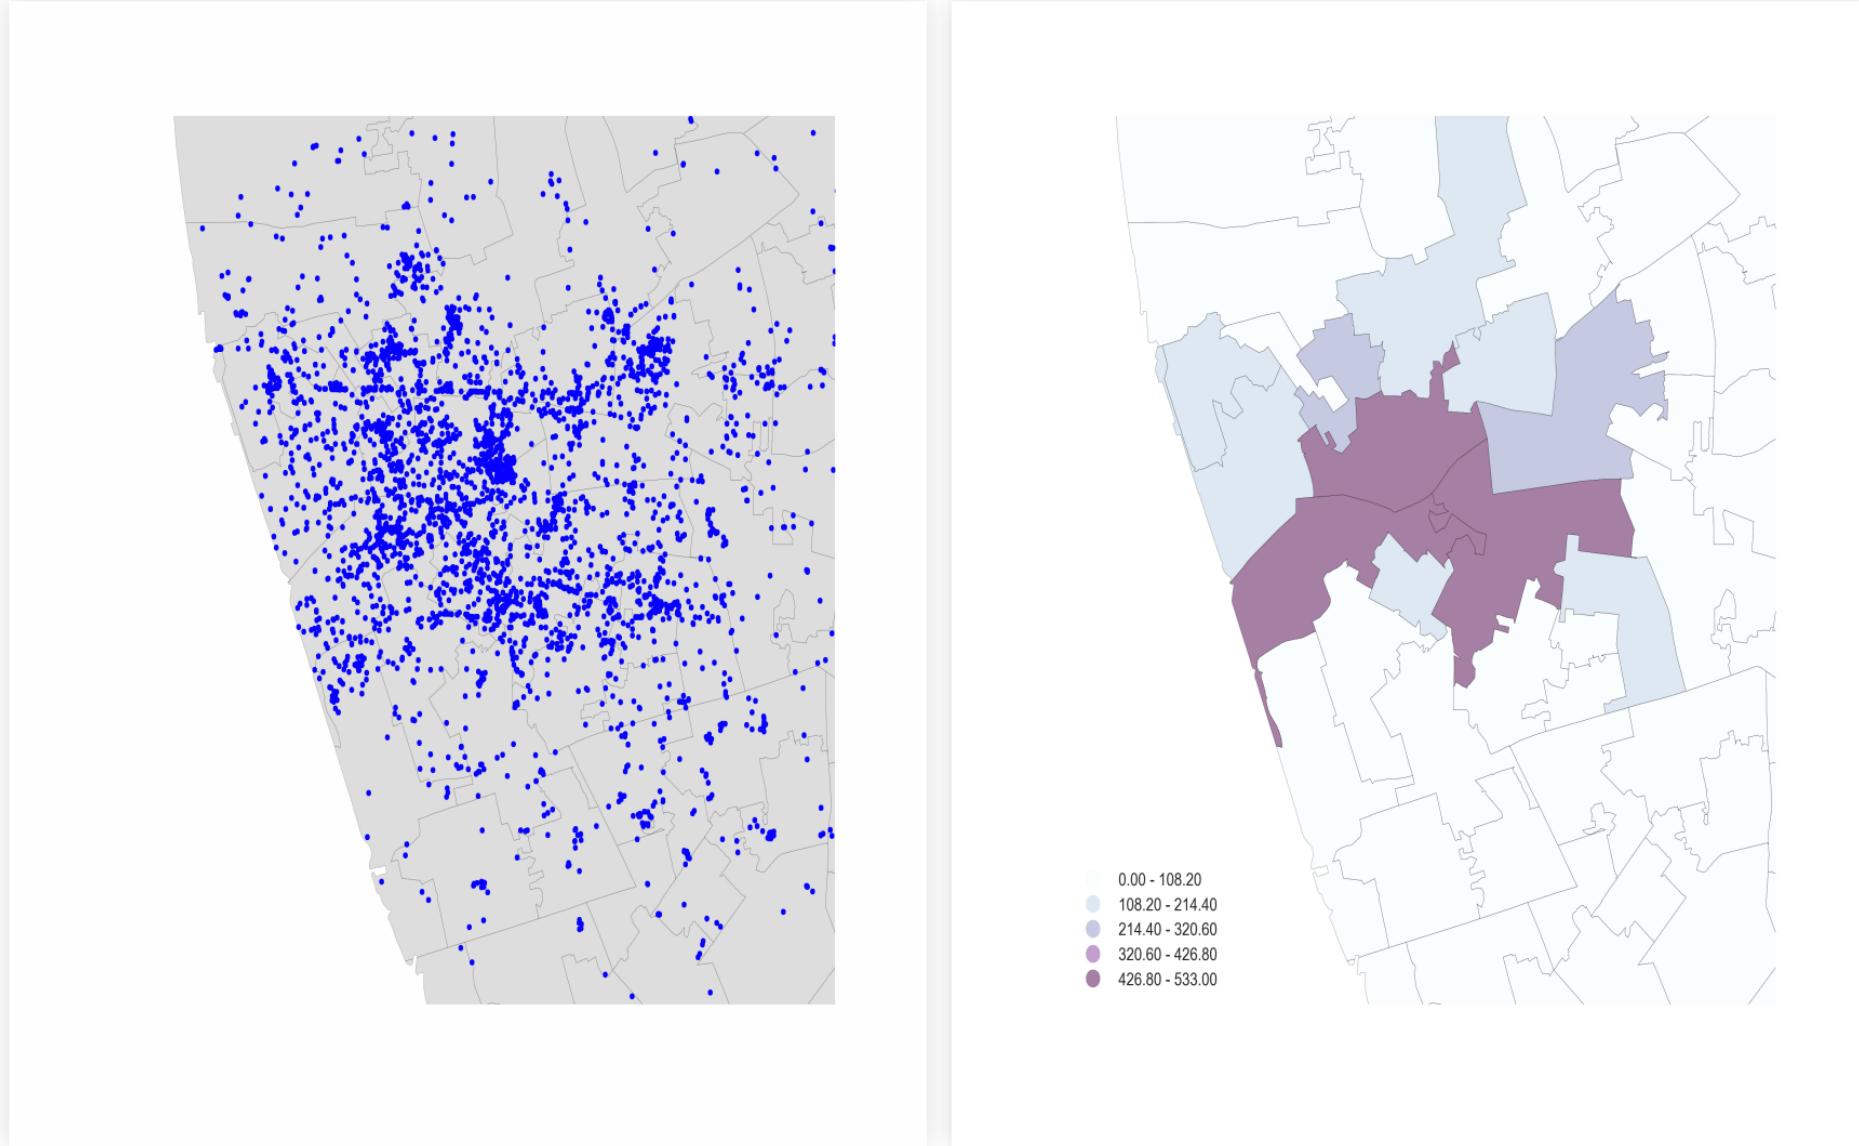
Aggregation

Points meet Polygons

Use polygon boundaries and count points per area

[Insert your skills for choropleth mapping here!!!]

But the polygons need to “make sense” (their delineation needs to relate to the point generating process).



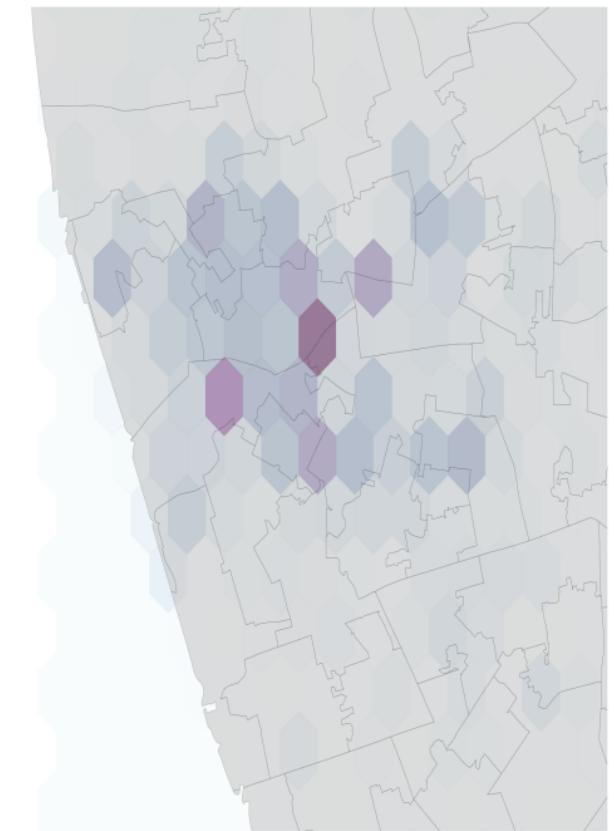
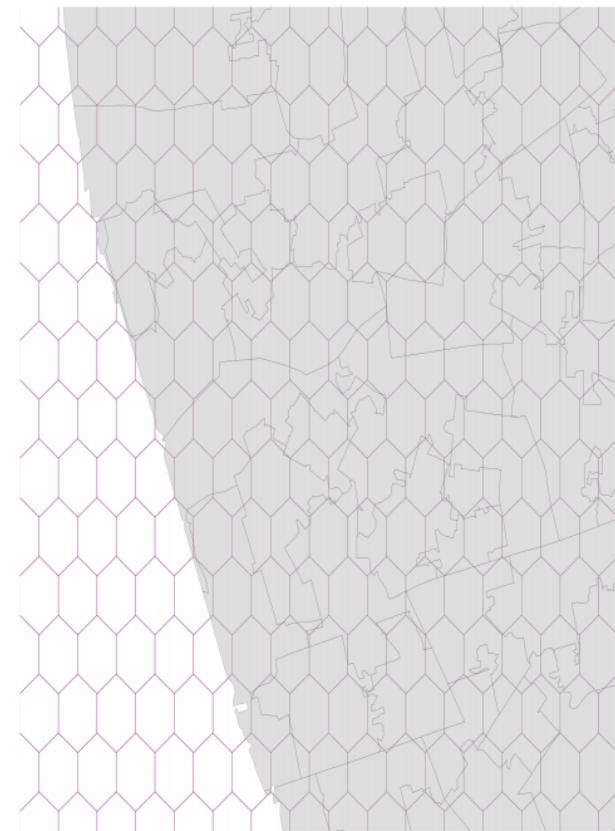
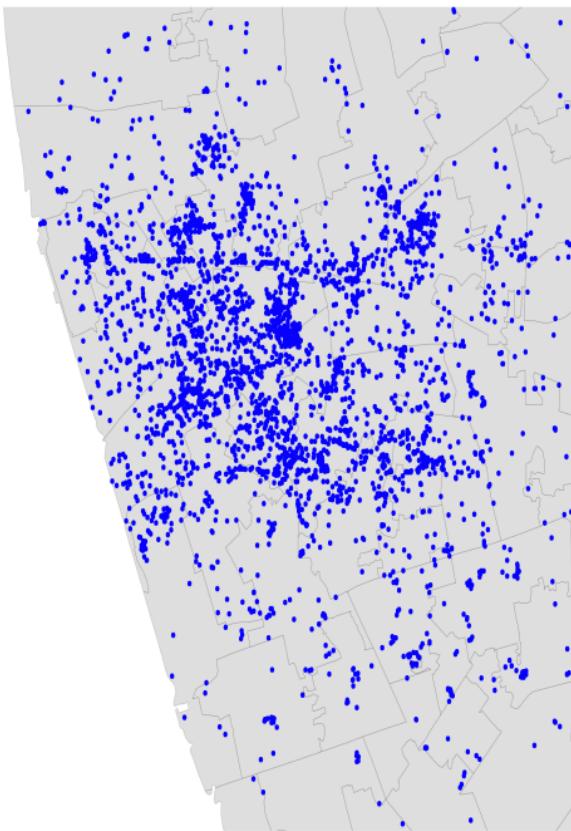
Hex-Binning

If no polygon boundary seems like a good candidate for aggregation...

...draw a **hexagonal** (or squared) tessellation!!!

Hexagons...

- Are regular (over census tracts)
- Exhaust the space (unlike circles)
- Have many sides (minimise boundary problems – think queen and rook!)



But...

(Arbitrary) aggregation may induce MAUP (see Lecture 6)

+

Points usually represent events that affect only part of the population and hence are best considered as rates (see Lecture 6)

Q: Which processes will you consider hexagon binning for?

... Points showing accessibility to amenities

- A. Yes
- B. No

Q: Which processes will you consider hexagon binning for?

... Points showing census locations

- A. Yes
- B. No

Q: Which processes will you consider hexagon binning for?

... Points showing cell-phone use in a location

- A. Yes
- B. No

Break



WATER



WALK



COFFEE OR TEA



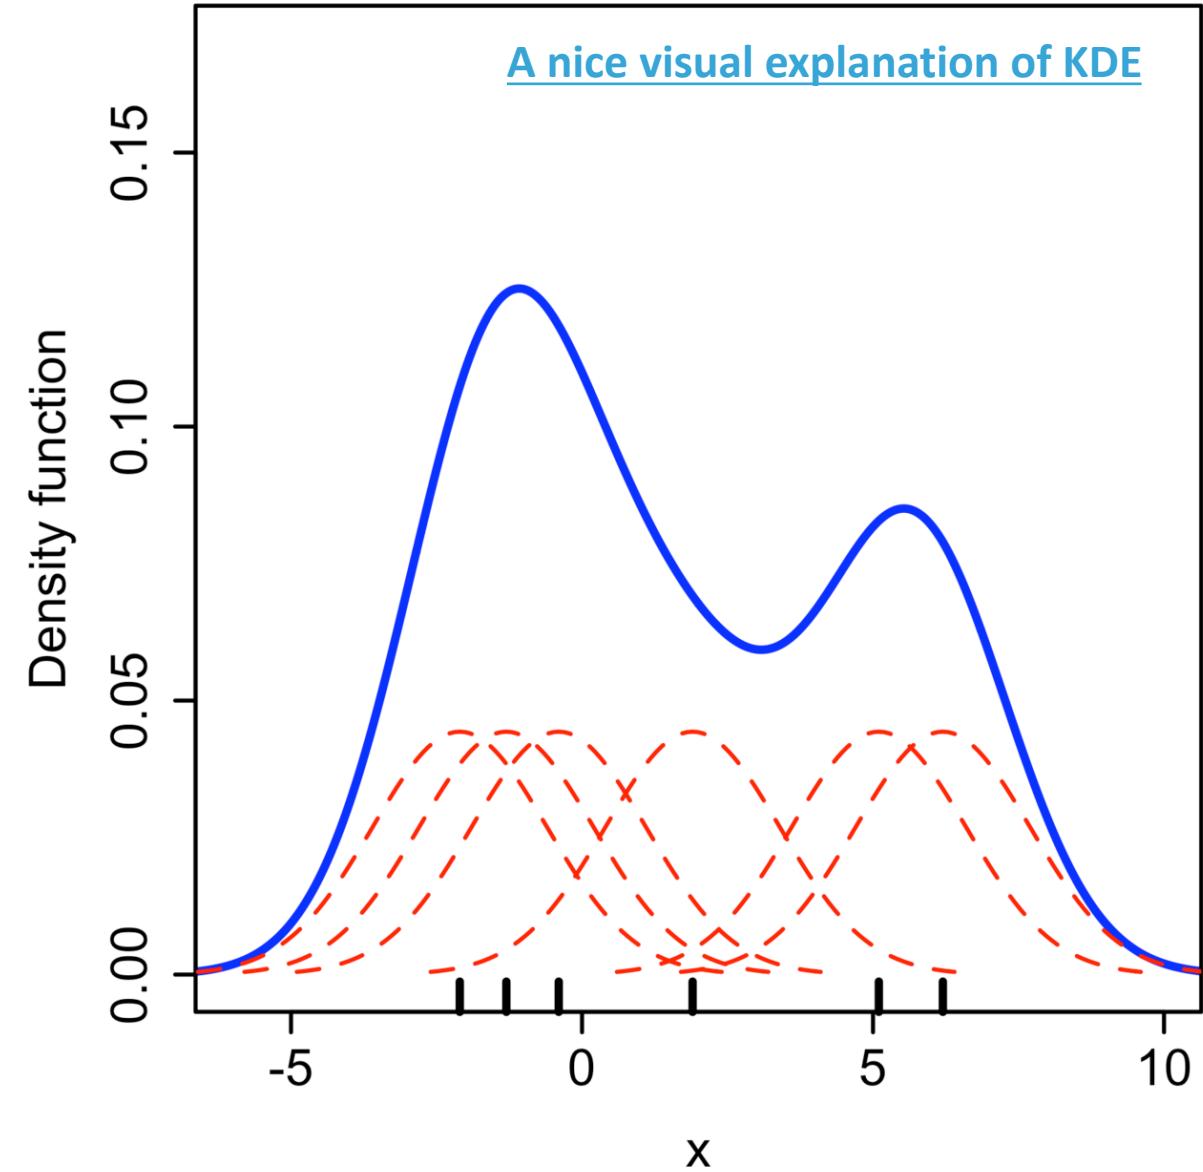
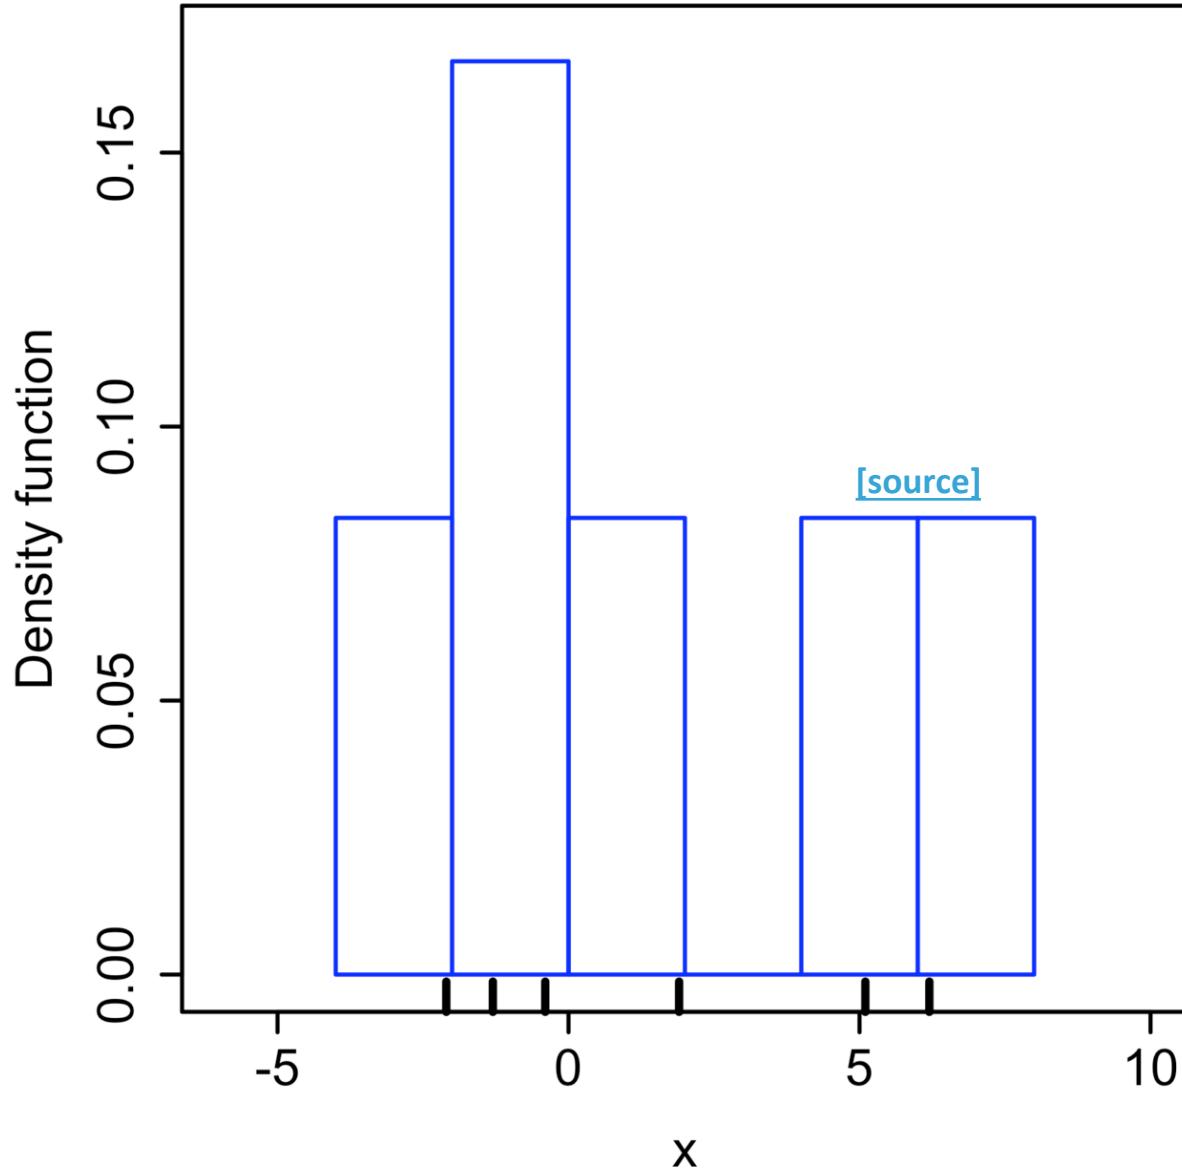
MAKE FRIENDS

Kernel Density Estimation

Kernel Density Estimation

Estimate the (*continuous*) observed *distribution* of a variable

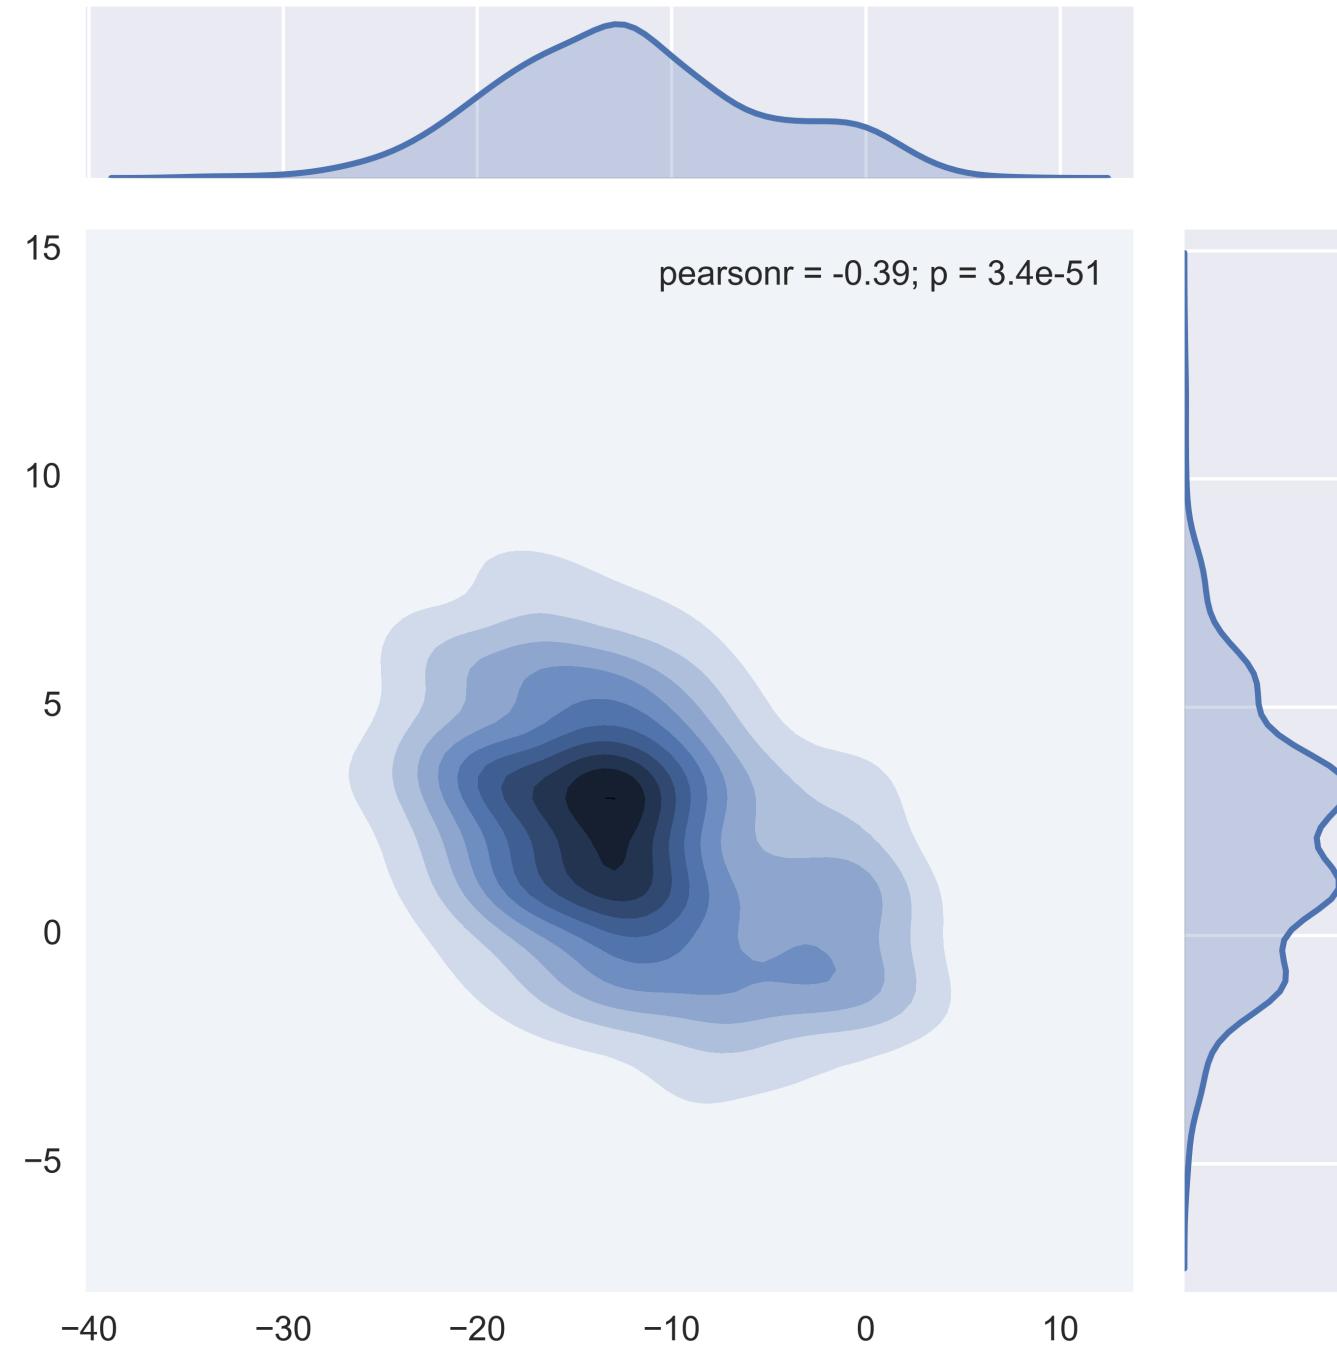
- Probability of finding an observation at a given point
- “Continuous histogram”
- Solves (much of) the MAUP problem, but not the underlying population issue

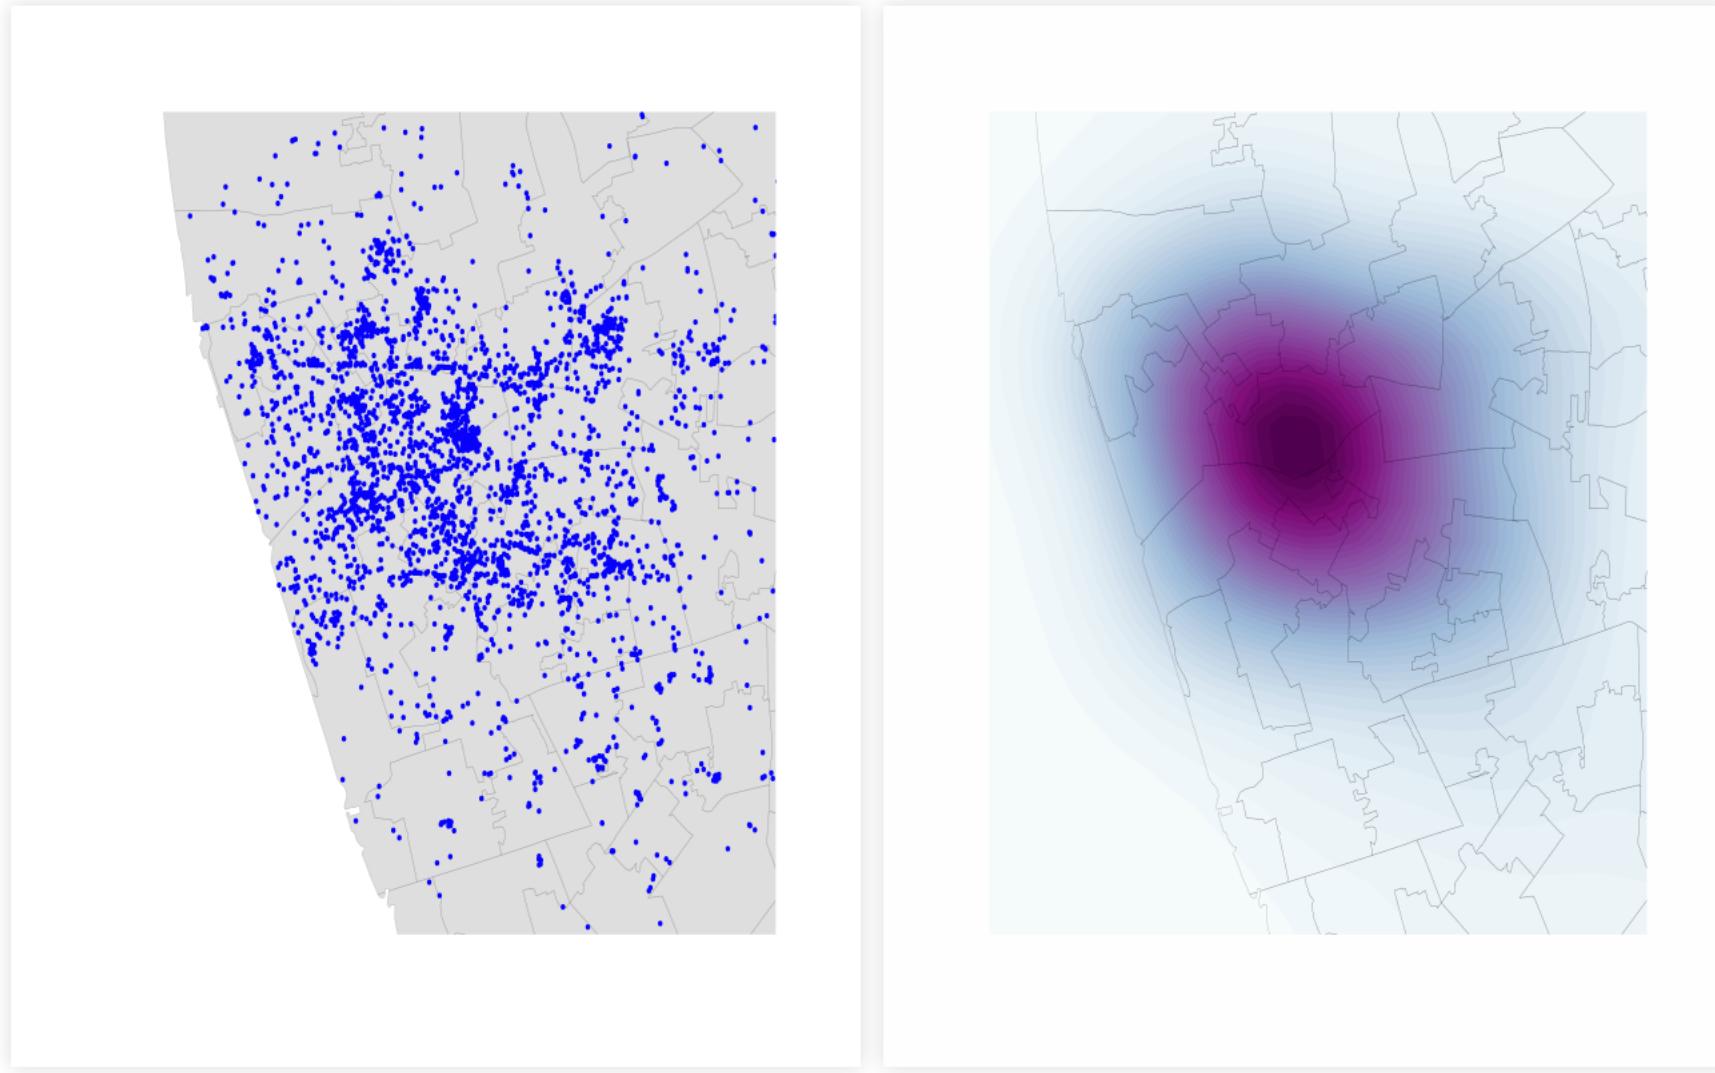


Bivariate (spatial) KDE

Probability of finding observations at a given point in space

- **Bivariate** version: distribution of **pairs of values**
- **In space**: values are coordinates (XY), locations
- Continuous “version” of a choropleth





Finding clusters of PPs

Concentrations/agglomerations of points over space, significantly more so than in the rest of the space considered

Huge literature spanning spatial analysis, statistics and computer science.

Today, we 'll look at [DBSCAN](#)

Q: When number of clusters (K) are known,
which algorithm do we use for clustering?

- A. DBSCAN
- B. K-Means
- C. K-Nearest Neighbours
- D. Support Vector Machines

Density
Based
Spatial
Clustering of
Applications with
Noise

When K is not known
↳ no. of clusters



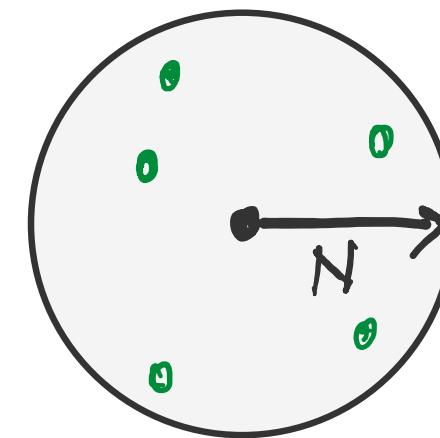
DBSCAN

- ↳ Set of points in space
- ↳ Neighbourhood N
- ↳ Density (minpts)

1. Set of points

Object → Space
(x, y)

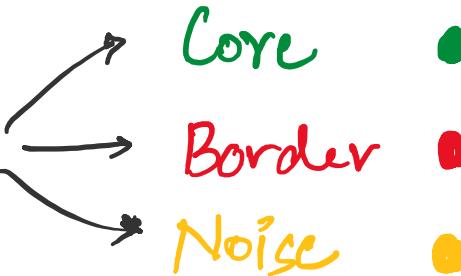
2. Neighbourhood



Steps

1. Label each pt. as

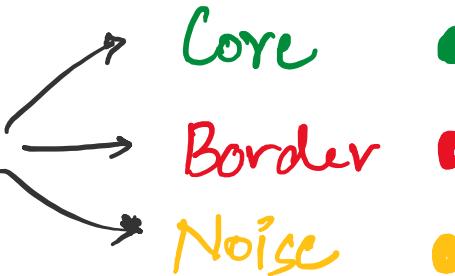
minpts = 4



Steps

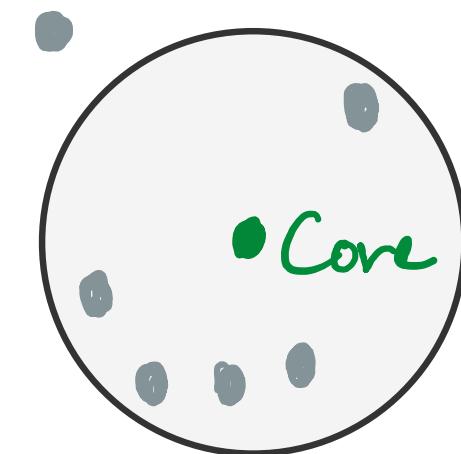
1. Label each pt. as

$$\underline{\text{minpts} = 4}$$



$$|N(x)| = 6$$

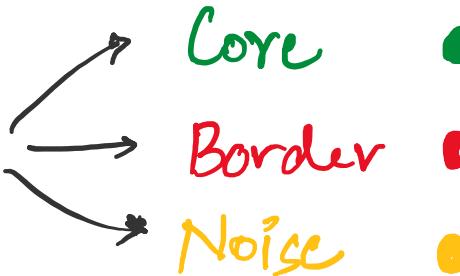
$$4 < 6$$



Steps

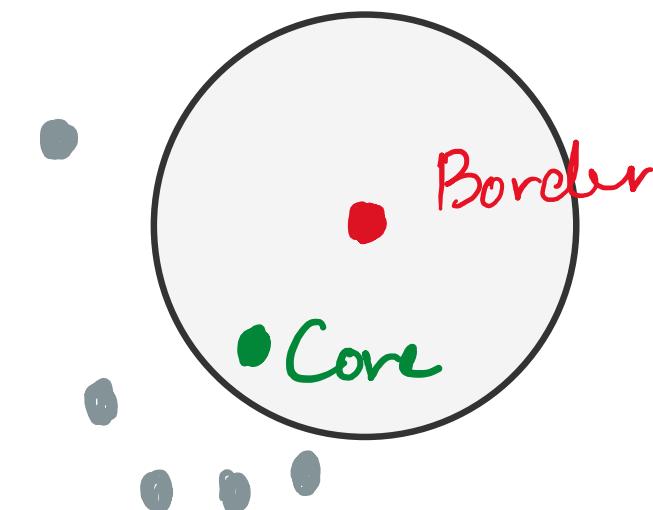
1. Label each pt. as

$$\underline{\text{minpts} = 4}$$



$$|N(x)|=2$$

$$4 > 2$$

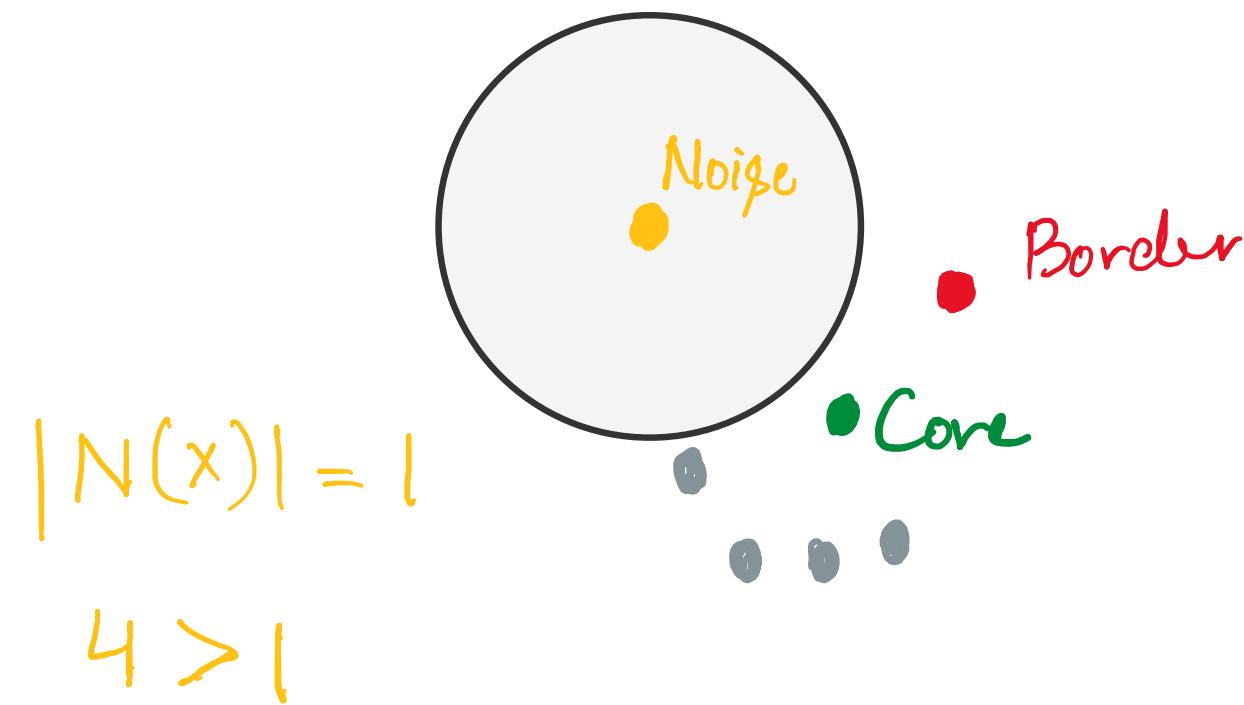


Steps

1. Label each pt. as

$$\underline{\text{minpts} = 4}$$

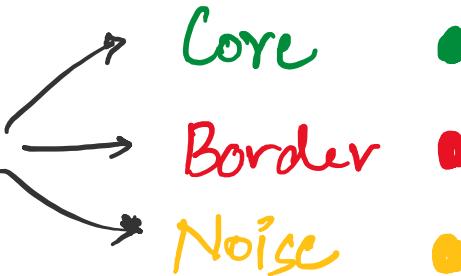
Core •
Border •
Noise •



Steps

1. Label each pt. as

$$\underline{\text{minpts}} = 4$$



2. Takes every **core** and performs a Depth First Search to find its neighbours

DBSCAN

Pros:

- Discover **any** number of clusters
- Clusters of varying **size** and **shape**
- Detect and **ignore outliers** in the data
- Not necessarily spatial
- Very fast to run so → scales relatively well → applicable to large datasets

Cons:

- Sensitive to **Neighbourhood** parameter
 - too small – sparse is noise
 - too large – dense merged together
- Not based on any probabilistic model (no inference)
- Hard to learn about the underlying process

For next class..



Finish Lab 08 to practice programming



No homework this week



Check Project – due in Week 10 on Friday at **2330**



See “To do before class” for every lecture (~ 1 hour of self study)



Read paper for **Discussion** session before every Friday



Post questions on the **Discussion** forum on Brightspace (especially on **Density Estimation** for this week)